

18 MARCH

01

Wk 09 • 060 Day

THURSDAY

# Applied Machine Learning (in python)

(Univ. of Michigan)

March 2018

Wk	M	T	W	T	F	S	S
09				1			03
10	5	6	7	8	9	10	11
11	12	13	14	15	16	17	18
12	19	20	21	22	23	24	25
13	26	27	28	29	30	31	

Scikit - Learn      | individualizar +  
SciPy  
NumPy  
Pandas  
matplotlib      cm.get\_cmap ('gnuplot')  
                      (small no. of features)  
                      → 3d → fixed 3D  
                      from ml\_toolkit import

K - Nearest Neighbour.

(for regression & classification)

↳ query point

↳ decision boundary.

Power parameter

→ Euclidean dist      (Minimise with  $P=2$ )

accuracy      .score()

Supervised Learning - (Label)

feature representation

Data instances / samples / examples ( $x$ )

Target Value ( $y$ )

	April 2010						
Wk	M	T	W	T	F	S	S
13	30	1	2	3	4	5	6
14	2	3	4	5	6	7	8
15	9	10	11	12	13	14	15
16	16	17	18	19	20	21	22
17	23	24	25	26	27	28	29

MARCH 18

WK 09 • 061 Day  
FRIDAY

09

## Classification

Discrete target

e.g. Yes/No

1/0

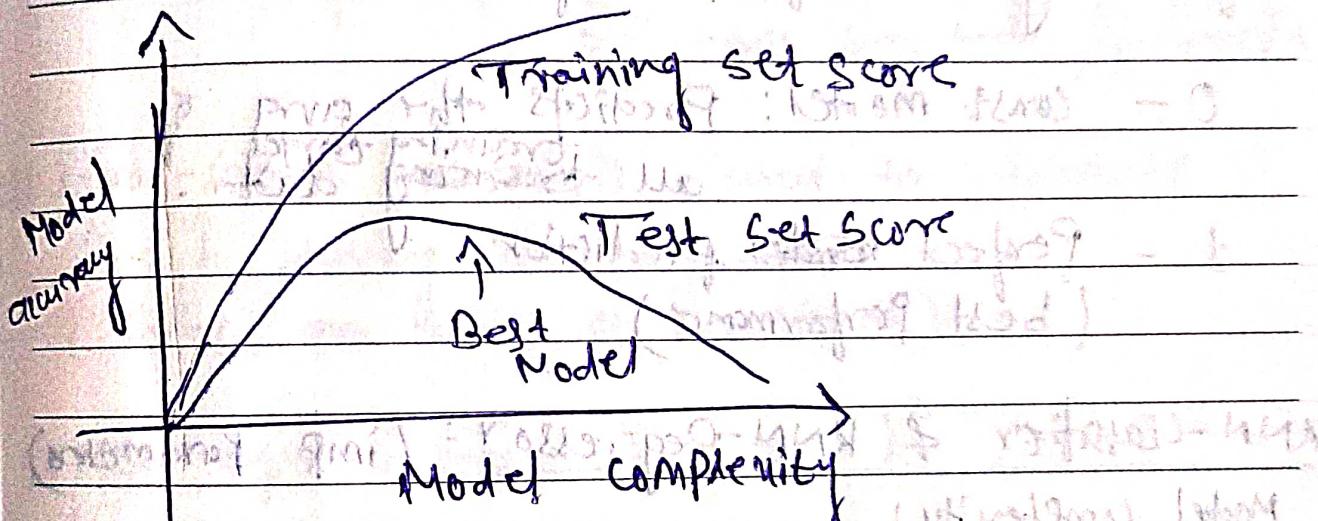
## Regression

continuous.

## Multiclass classifier

(More discrete  
values)

## Multi-label classification



## Overfitting & Underfitting

for KNN  $\rightarrow k=1 \rightarrow$  overfitting.

$k=5 \rightarrow$  good

$k=10 \rightarrow$  underfitting

for the last example.

$k \uparrow \rightarrow$  Risk of overfitting  $\uparrow$  for KNN

18 MARCH

03

Wk 09 • Day 2

SATURDAY

March 2018						
Wk	M	T	W	T	F	S
09				1	2	3
10	5	6	7	8	9	10
11	12	13	14	15	16	17
12	19	20	21	22	23	24
13	26	27	28	29	30	31

KNN can be used for both classification & regression.

#  $R^2$  - regression score < 0,1  
(co-efficient of determination)  
→ Measures how well the prediction model for regression fits the given data.

0 - const model: Predicts the avg of all training targets

1 - Perfect predictor.  
(best performance)

KNN-Classifier & KNN-Regression: (imp Parameters)

Model Complexity:

n\_neighbours : default = 5.

04 Sunday

Model-fitting:

metric: dist from best data points.

default = P=2.

April 2010

S	M	T	W	T	F	S	S
1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22	23	24
25	26	27	28	29	30	1	2

MARCH '18

Wk 10 • 064 Day  
MONDAY

05

Least-Square : L.P.Input instance (feature vector)  $\vec{x} = (x_0, x_1, \dots, x_n)$ Predicted output:  $\hat{y} = \hat{w}_0 x_0 + \hat{w}_1 x_1 + \dots + \hat{w}_n x_n + \hat{b}$ 

Parameters to estimate:  $\hat{w} = (\hat{w}_0, \dots, \hat{w}_n)$ : feature weights / model coefficients.  
 $\hat{b}$ : const bias term / intercept.

# Least-Square (method) is used to estimate  $\hat{w}$  &  $\hat{b}$  during training.  
 (There are also other methods)

mean square Error =  $(y_i - \bar{y})^2$ 

In Linear Regression, there is no parameter to control Model complexity.

obj func - optimizr

loss func - minimize

$$RSS(w, b) = \sum_{i=1}^n (y_i - \hat{y})^2$$

(Sum of Squared Diff)  
(RSS)

(Residual Sum of Squares)

(RSS)

'18 MARCH

06

Wk 10 • 065 Day

TUESDAY

March 2018

Wk	M	T	W	T	F	S	S
09				1	2	3	4
10	5	6	7	8	9	10	11
11	12	13	14	15	16	17	18
12	19	20	21	22	23	24	25
13	26	27	28	29	30	31	

sklearn : Linear Regression

$w_{\text{linreg}} = \text{Linreg. Coef.}$

$b_{\text{linreg}} = \text{Linreg. intercept}$

((underscore denotes a quantity derived from training data))

## LR: Ridge, Lasso & Polynomial Regression

$$RSS_{\text{Ridge}}(w, b) = \sum_{i=1}^N (y_i - (w \cdot x_i + b))^2 + \alpha \sum_{j=1}^P w_j^2$$

Regularization

↓

It prevents overfitting by: restricting model, (reducing its complexity)

L2 Regularization:  $\min \sum \text{of squares of } w_j \text{ of entries.}$

$\alpha \uparrow$  regularization  $\uparrow \Rightarrow$  more simple model.

$\alpha = 1$  (default)

# from sklearn.linear\_model import Ridge,

# feature pre-processing & normalization.

	April 2018						
Mo	Tu	We	Th	Fr	Sa	Su	
1					1		
2					2		
3		3	4	5	6	7	8
4		9	10	11	12	13	14
5		16	17	18	19	20	21
6		23	24	25	26	27	28
7							29

MARCH '18  
07  
WEDNESDAY

minMaxScaling.

$[0, 1] \rightarrow \text{Value}$ .

$$m'_i = \frac{x_i - x_{i, \text{Min}}}{x_{i, \text{Max}} - x_{i, \text{Min}}}$$

# from sklearn.preprocessing import MinMaxScaler  
(using transform())

\* Data Leaking \*.

Small no. of training data compared to no. of features  $\Rightarrow$  Regularization works well.

## Lasso Regn

$$\text{RSS} = \dots + \lambda \sum_{j=1}^p |w_j|$$

L1 Penalty

Sparse solution: kind of feature selection.

# from

import Lasso

18 MARCH

08

Wk 10 • 057 Day

THURSDAY

March 2018

Wk	M	T	W	T	F	S	SU
09				1	2	3	4
10	5	6	7	8	9	10	11
11	12	13	14	15	16	17	18
12	19	20	21	22	23	24	25
13	26	27	28	29	30	31	

Polynomial feature Transformations

$$x = (n_0, n_1) \rightarrow w^T = (n_0, n_1, n_0^2, n_0 n_1, n_1^2)$$

Non-linear basis functions.

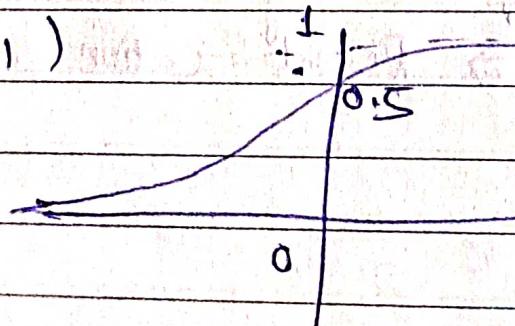
Logistic Regression :-

(Used for Classification)

$$\hat{y} = \text{logistic}(\hat{b} + \hat{w}_1 \cdot n_1 + \dots + \hat{w}_n \cdot n_n)$$

$$= \frac{1}{1 + \exp[-(\hat{b} + \hat{w}_1 \cdot n_1 + \dots)]}$$

$$y \in (0, 1)$$



# from sklearn.linear\_model import LogisticRegression

✓ regularization is 'on' by default.

✓ Parameter C controls amount of regularization.

(↑ Regularization ↓) (default 1.0)

April 2018

M		T	W	T	F	S
WK	M	T	W	T	F	S
13	30	1	2	3	4	5
14	2	3	4	5	6	7
15	9	10	11	12	13	14
16	16	17	18	19	20	21
17	23	24	25	26	27	28

MARCH '18

WK 10 • 068 Day  
FRIDAY

09

Linear classifiers! SVMs,

$$x \rightarrow \boxed{f(x)} \rightarrow y$$

$$f(w, w_1, b) = \text{sign}(w \cdot x + b)$$

dot product:

$$(w_1, w_2) \cdot (v_1, v_2)$$

$$= w_1 v_1 + w_2 v_2$$

→ Classifier Margin:

Margin width → the decision boundary area can be increased before hitting a data point.

→ Margins in classifier?

Linear classifier with margin is a Linear SVM. (LSVM)

# from sklearn.svm import SVC.

\* Larger C: less regularization.

→ fit the training data as well as possible

→ each individual data point is imp to classify correctly.

\* Smaller C: More regularization!

→ More tolerant of errors on individual data points,

'18 MARCH

10

Wk 10 • 069 Day

SATURDAY

March 2018

Wk	M	T	W	T	F	S
09				1	2	3
10	5	6	7	8	9	10
11	12	13	14	15	16	17
12	19	20	21	22	23	24
13	26	27	28	29	30	31

Multiclass classification

Kernelized Support Vector Machine - (SVM)

(both classification & Regression)

→ RBF = Radial Basis Function Kernel.

→ Polynomial Kernel.

RBF

$$k(x, x') = \exp[-\gamma \|x - x'\|^2]$$

↑  
gamma ( $\gamma$ ) : kernel width parameter.

{ Kernel type.  
 $\gamma$  value  
 11 Sunday C - parameter. } → These 3 affect kernel complexity.

Cross Validation:

$k$ -fold cross validation

# from sklearn.model\_selection

import

cross\_val\_score

							April 2018
Wk	M	T	W	T	F	S	S
13	30	1	2	3	4	5	6
14	2	3	4	5	6	7	8
15	9	10	11	12	13	14	15
16	16	17	18	19	20	21	22
17	23	24	25	26	27	28	29

MARCH '18

Wk 11 • 071 Day  
MONDAY

12

stratified k-fold Cross Validation.

Leave one out cross validation!  $k=1$  data for test set.

validation curve

## Decision Trees

(Supervised). Standard method after linear.

If then conditions = series of yes/no questions.

Root Node, Leaf Node.

# from sklearn.tree import DecisionTreeClassifier

Controlling the Model Complexity:

Parameters  $\Rightarrow$  max\_depth

$\hookrightarrow$  Max # of leaf nodes = max\_leaf\_nodes.

$\hookrightarrow$  Min sample to

consider splitting = min\_samples\_leaf

18 MARCH

13

Wk 11 • 079 Day  
TUESDAY

Re-Read : SVM

Ridge Regression

Lasso

Regression

March 2018						
Wk	M	T	W	T	F	S
09	5.	6	7	8	9	10
11	12	13	14	15	16	17
12	19	20	21	22	23	24
13	26	27	28	29	30	31

feature importance :

(How imp is a feature to overall prediction accuracy.)

f.i=0  $\Leftrightarrow$  not used in prediction

1  $\Leftrightarrow$  the feature predicts the target perfectly

## Model Evaluation & Selection

Accuracy also has drawbacks.

Unbalanced Clab.

Dummy Classifier:

$\hookrightarrow$  Provide a null metric (e.g., null accuracy) baseline.

$\hookrightarrow$  Dummy Classifier should not be used for real problems.

~~$\checkmark$  Accuracy achieved by always picking the most frequent class.~~

$\hookrightarrow$  It ignores the input data.

Common settings for the "Strategy" parameter for DummyClassifier in sklearn:

$\hookrightarrow$  Most-frequent: Predicts most freq. label.

$\hookrightarrow$  Stratified: random pred.

MARCH 18

14

Wk	M	T	W	T	F	S	S
13	30	1	2	3	4	5	6
14			7	8	9	10	11
15			12	13	14	15	16
16			18	19	20	21	22
17			23	24	25	26	27

Wk 11 • 073 Day

WEDNESDAY

↳ Uniform: generated Preqd uniformly at random.

↳ Const: Always Preqd a const label provided by user.

AUC (Area Under the Curve) :

Confusion matrix :

# from sklearn.metrics import

confusion\_matrix

True		False			
True		False			
-		+			
					← Confusion matrix
False	True				
-ve	+ve				

false +ve = Type-1 error

false -ve = Type-2 error

Confusion Matrices & Basis Evaluation Metrics:

$$\text{Classification Error} = \frac{\text{# Errors (FP+FN)}}{\text{# total (FP+FN+TP+TN)}}$$

(1 - Accuracy)

$$\text{Recall} = \frac{TP}{TP+FN}$$

(frac of all +ve instances does the classifier correctly identify as +ve)

18 MARCH

15

WX 11 • 074 DAY

THURSDAY

March 2018						
Wk	M	T	W	T	F	S
09				1	2	3
10	5	6	7	8	9	4
11	12	13	14	15	16	10
12	19	20	21	22	23	17
13	26	27	28	29	30	24
						25
						31

\* Precision : fraction of the predictions that were correct?  $\frac{TP}{TP+FP}$

\* False Positive Rate (FPR) :  $\frac{FP}{TP+FP}$  (or Specificity)

# F1-Score = AIM of Precision & Recall.  

$$F1 = 2 \cdot \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

# F1 is a case of F-score.

$$F\text{-Score} = \frac{1 + \beta^2}{(1 + \beta^2) \cdot \frac{\text{Precision} \times \text{Recall}}{(\beta^2 \times \text{Precision}) + \text{Recall}}}$$

# from sklearn.metrics import accuracy\_score, precision\_score, recall\_score, f1\_score, classification\_report:

Classify Decision functions:

decision\_function(),

Predict\_Proba().

	M	T	W	T	F	S	S
13	30		4	5	6	7	8
14	2	3	10	11	12	13	14
15	9	10	11	12	13	14	15
16	16	17	18	19	20	21	22
17	23	24	25	26	27	28	29

16

Precision Recall & ROC curves:

Receivers Operating Characteristics Curves

Evaluation  
Multi-class classification!

micro vs Macro Average.

(each class has equal weight)

Each instance has equal wt.

Largest classes have most influence.

↳ If classes have same no. of instances,  
macro  $\approx$  micro

↳ If some classes have much larger(<sup>more</sup>, instances)  
than others =

wt metric towards largest one - micro

wt 1 " " smallest " - macro

↳ micro << macro, larger class has poor metric performance

↳ micro >> macro, smaller " "

'18 MARCH

17

Wk 11 • 076 Day

SATURDAY

March 2018

Wk	M	T	W	T	F	S
09				1	2	3
10	5	6	7	8	9	4
11	12	13	14	15	16	11
12	19	20	21	22	23	18
13	26	27	28	29	30	25
						31

Regression Evaluation

Optimizing Classifiers for different Eval. metrics

data testing :-

Dimensionality Red<sup>↑</sup> & Manifold Learnings

PCA: Principal Component Analysis

```
# from sklearn.decomposition import PCA
```

MDS - Multidimensional Scaling,

```
# from sklearn.manifold import MDS
```

Clustering

18 Sunday

	April 2018						
M	W	T	W	T	F	S	S
13	14	15	16	17	18	19	20
19	20	21	22	23	24	25	26
27	28	29					

MARCH 18

Wk 19 • 078 Day

MONDAY

19

## Naïve Bayes Classifier

3 types (available in sklearn)

↳ Bernoulli : binary features (e.g. word present also)

↳ Multinomial : discrete features (e.g. word counts)

↳ gaussian : Continuous / real valued features.

↳ # from sklearn.naive\_bayes import GaussianNB

supports partial fit.

## Random Forest

→ An ensemble of trees.

sklearn.ensemble module:

classifier : RandomForestClassifier

Regression : RandomForestRegressor

n\_estimators Parameter.

Bootstrap Samples

max\_features

18 MARCH

20

Wk 12 • 079 Day

TUESDAY

March 2018

Wk	M	T	W	T	F	S	S
09				1	2	3	4
10	5	6	7	8	9	10	11
11	12	13	14	15	16	17	18
12	19	20	21	22	23	24	25
13	26	27	28	29	30	31	

gradient boosted decision trees

(weak learners)

(learning rate)

(Boosting) how gradient strength (improves) Neural Networks

Multi Layer Perceptron

(feed forward NN)

(tanh activation funcn)

Import MLPRegressor from sklearn.neural\_network

IMP parameters:

↳ hidden\_layer\_sizes

↳ alpha

↳ activation

Data Leakage